

Assessment for AI Security Researcher

LLM Security Evaluation

Scope

You are provided a set of malicious prompts. Your task is to:

- Determine whether the LLM is vulnerability to attack based on it's response to the prompt.
- Run controlled experiments against an AWS Bedrock-hosted model or any cloud hosted LLM model.
- Determine whether the model **complies** with the malicious intent ("falls for it") or **refuses/safeguards**.
- Produce a report with the findings.

Steps

1) Build an web based evaluation harness

Create a web interface which allows the prompts to be uploaded. And provide option to download the reports into an appropriate file format.

2) Prompt analysis & taxonomy

For each prompt, classify it **before** running it against the model:

- **Benign**: normal user request without harmful intent.
 - **Malicious**: attempts to bypass safety, security, exfiltrate sensitive info.. etc
 - The evaluation must be done automatically.
-

3) Model behavior testing ("does it fall for it?")

Run the prompts through the chosen Bedrock model or any online LLM. You may choose any LLM of any version. The purpose is not to find vulnerabilities in latest LLM.

For each prompt tested, assess:

- Does the response **comply** with the malicious intent?
- Does it **refuse** appropriately?
- Any **leakage** of hidden instructions, sensitive data, or unsafe content?
- The evaluation must be done automatically.

Deliverable: Web interface demonstration and testing report.